


MID_SIZED RETAIL COMPANY DIMENSIONAL MODELLING REPORT

ACSP - Senior Data Analyst - Stage Eight

INTRODUCTION

Tools employed
Relevant links



Conclusion and Relevance

Assignment Requirements

1. Data Profiling & Inference
2. Dimensional Modeling Strategy
3. Slowly Changing Dimensions (SCD)
4. ETL/ELT Strategy
5. Advanced Analysis & Scalability

INTRODUCTION

- ❖ The dimensional modelling was carried out per client instructions given on dorian-twister-cea.notion.site
- ❖ The Dimension was done was carried out in PGAdmin using SQL script. Relevant screenshot are attached to the report.
- ❖ PDadmin was used to create database function to create a database called storemanagement
- ❖ [Link to the dimension Model](#)
- ❖ [Link to the SQL codes for the database creation and Triggers and SCD.](#)

1. Data Profiling and Inference

Identifying Dimensions in the Invoice Dataset

❖ Overview

- Dimensions provide descriptive context to transactional data
- They help categorize, filter, and analyze key business entities.

Identified Dimension Tables: Four dimensions were identified and they are given below

❖ Customer Dimension (customer_dim):

- Captures details of customers involved in transactions.
- Attributes:
 - customer_id (Primary Key)
 - Customer_name
 - Address
 - Telephone
- The code in the screenshot was used to load this dimension to the postgres database using the query tool.

```
-- Create customer dimension table  
CREATE TABLE customer_dim (  
    customer_id VARCHAR(50) PRIMARY KEY,  
    customer_name TEXT NOT NULL,  
    address TEXT,  
    telephone VARCHAR(20)  
);
```

1. Data Profiling and Inference

Identified Dimension Tables

❖ Supplier Dimension (supplier_dim)

- When goods are out of stock or in the re-order levels, retailers reach out to their suppliers to replenish stocks.
- This stores information about suppliers providing goods.
- **Attributes:**
 - supplier_id (Primary Key)
 - Supplier_name
 - Address
 - Telephone

❖ Store Dimension (store_dim)

- This dimension store information of physical or online stores where sales occur
- **Attributes:**
 - store_id (Primary Key)

```
-- Create supplier dimension table
CREATE TABLE supplier_dim (
    supplier_id VARCHAR(50) PRIMARY KEY,
    supplier_name TEXT NOT NULL,
    address TEXT,
    telephone VARCHAR(20)
);

-- Create store dimension table
CREATE TABLE store_dim (
    store_id VARCHAR(50) PRIMARY KEY,
    store_name TEXT NOT NULL,
    store_location TEXT,
    store_contact VARCHAR(20)
);
```

1. Data Profiling and Inference

Identified Dimension Tables

- Store_name
- Store_location
- Store_contact

❖ Product Dimension (product_dim)

- It contains details about products sold in transactions
- Attributes:
 - product_id (Primary Key)
 - product_name
 - quantity_available
 - Unit_price
 - Cost_price

❖ The data types of the attributes of each dimensions are given in each code snippet

```
-- Create product dimension table
CREATE TABLE product_dim (
    product_id VARCHAR(50) PRIMARY KEY,
    product_name TEXT NOT NULL,
    quantity_available INT DEFAULT 0,
    unit_price DECIMAL(10,2) NOT NULL,
    cost_price DECIMAL(10,2) NOT NULL
);
```

1. Data Profiling and Inference

Determining Fact Table Granularity

❖ Overview

- Granularity defines the level of detail stored in the fact table.
- It impacts reporting, performance, and the depth of analysis.

Established Fact Table Granularity: Two Fact table were established and they are given below

❖ Purchase Invoice Fact Table (**purchase_invoice_fact**):

- **Grain:** Each record represents a single product purchase per invoice line
- **Justification:**
 - Captures itemized purchases for cost tracking.
 - Supports supplier performance analysis.
 - Enables procurement trend analysis.

❖ Sales Invoice Fact Table (**sales_invoice_fact**):

- **Grain:** Each record represents a single product sale per invoice line.

1. Data Profiling and Inference

Established Fact Table Granularity

- ❖ **Sales Invoice Fact Table (sales_invoice_fact):**
 - Allows detailed sales trend analysis.
 - Supports customer purchase behavior insights.
 - Enables store-wise performance evaluation.

- ❖ **Key Business Insights Enabled:**
 - **Profitability Analysis:** Compare unit prices vs. cost prices at an item level.
 - **Inventory Management:** Track product movement across stores.
 - **Customer & Supplier Performance:** Identify top buyers and best-performing suppliers.
 - **Trend Forecasting:** Analyze demand fluctuations over time

- ❖ The code snippet containing the dtypes, attributes and other relevant information of the fact table are given in the next slide

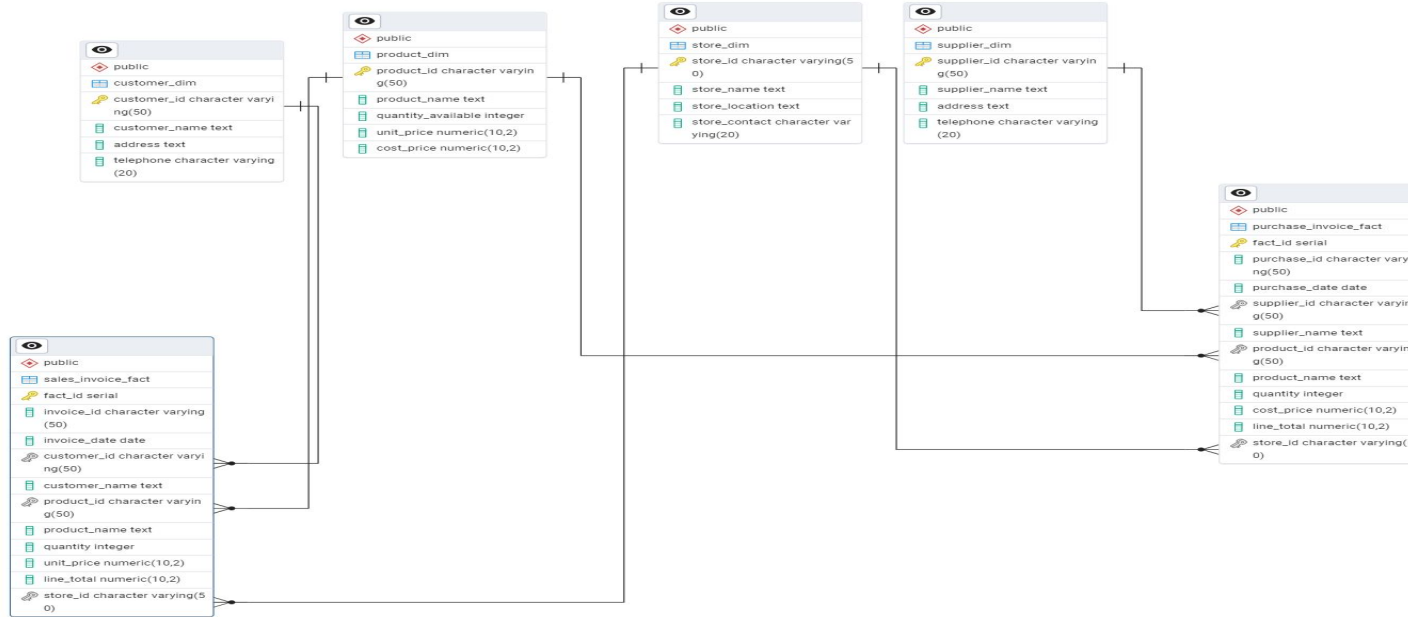
1. Data Profiling and Inference

```
-- Create purchase invoice fact table with a surrogate primary key.
CREATE TABLE purchase_invoice_fact (
    fact_id SERIAL PRIMARY KEY,          -- Surrogate key for unique identification
    purchase_id VARCHAR(50),             -- Now not unique; can be repeated
    purchase_date DATE NOT NULL,
    supplier_id VARCHAR(50) REFERENCES supplier_dim(supplier_id) ON DELETE CASCADE,
    supplier_name TEXT NOT NULL,
    product_id VARCHAR(50) REFERENCES product_dim(product_id) ON DELETE CASCADE,
    product_name TEXT NOT NULL,
    quantity INT NOT NULL CHECK (quantity >= 0),
    cost_price DECIMAL(10,2) NOT NULL CHECK (cost_price >= 0),
    line_total DECIMAL(10,2) GENERATED ALWAYS AS (COALESCE(quantity * cost_price, 0)) STORED,
    store_id VARCHAR(50) REFERENCES store_dim(store_id) ON DELETE CASCADE
);

-- Create sales invoice fact table with a surrogate primary key.
CREATE TABLE sales_invoice_fact (
    fact_id SERIAL PRIMARY KEY,          -- Surrogate key for unique identification
    invoice_id VARCHAR(50),             -- Now not unique; can be repeated
    invoice_date DATE NOT NULL,
    customer_id VARCHAR(50) REFERENCES customer_dim(customer_id) ON DELETE CASCADE,
    customer_name TEXT NOT NULL,
    product_id VARCHAR(50) REFERENCES product_dim(product_id) ON DELETE CASCADE,
    product_name TEXT,
    quantity INT NOT NULL CHECK (quantity >= 0),
    unit_price DECIMAL(10,2) NOT NULL CHECK (unit_price >= 0),
    line_total DECIMAL(10,2) GENERATED ALWAYS AS (COALESCE(quantity * unit_price, 0)) STORED,
    store_id VARCHAR(50) REFERENCES store_dim(store_id) ON DELETE CASCADE
);
```

2. Dimensional Modeling Strategy

❖ Star Schema design



2. Dimensional Modeling Strategy

❖ Key Benefits of The Star Schema design

- Fast Query Performance – Optimized for aggregations and analytics
- Simplified Reporting – Enhances sales, inventory, and supplier insights.
- Scalability – Easily extends by adding new dimensions or measures.
- Efficient Storage – Reduces redundancy while maintaining data integrity.

Handling Degenerate Dimensions

- ❖ Degenerate dimensions are identifiers that exist in fact tables but have no corresponding dimension table.

Identified Degenerate Dimensions

- ❖ **purchase_id (in purchase_invoice_fact)**
 - Represents a unique purchase transaction.
 - No additional descriptive attributes justify a separate dimension table
- ❖ **invoice_id (in sales_invoice_fact)**
 - Represents a unique sales transaction.
 - Acts as a reference for order-level aggregation and tracking.
- ❖ **Management Strategy**
 - **Stored in the Fact Tables** – Since they have no descriptive attributes, they remain in the respective fact tables.
 - **Optimized for Query Performance** – Indexed for efficient filtering and lookups.

2. Dimensional Modeling Strategy

❖ **Management Strategy**

- **Stored in the Fact Tables** – Since they have no descriptive attributes, they remain in the respective fact tables.
- **Optimized for Query Performance** – Indexed for efficient filtering and lookups.
- **Facilitates Drill-Down Analysis** – Enables tracking of individual transactions.
- **Supports Data Integrity** – Ensures accurate reconciliation between transactions.

❖ **Business Insights Enabled**

- **Sales Trend Analysis** – Group sales data by invoice_id to analyze transaction patterns.
- **Purchase Order Tracking** – Aggregate purchases using purchase_id for procurement efficiency.
- **Operational Auditing** – Helps track specific transactions for validation and troubleshooting.

3. Slowly Changing Dimensions (SCD)

❖ Understanding SCD

- SCD refers to handling historical changes in dimension attributes over time.
- Common SCD types: Type 1 (Overwrite), Type 2 (Versioning), Type 3 (Limited History).

❖ Chosen Approach: SCD Type 2

- Tracks Historical Changes – Maintains multiple versions of a record.
- Adds Surrogate Keys – Ensures unique historical records.
- Uses Effective & Expiration Dates – Identifies active vs. archived records.
- Current Flag Indicator – Marks the latest version of the record.

❖ Implementation Details

- Customer Dimension (customer_dim_scd)
 - Tracks customer details like name, address, and telephone over time.
 - Trigger (trg_scd_customer_dim_update) archives old records before an update.
- Product Dimension (product_dim_scd)
 - Maintains historical product attributes like price and quantity available.
 - Trigger (trg_capture_product_history) updates expiration dates of old records.

3. Slowly Changing Dimensions (SCD)

❖ Trade-offs & Considerations

- Pros
 - Enables historical trend analysis and auditing.
 - Preserves a full history of changes for business intelligence.
- Cons
 - Increased storage due to multiple versions.
 - More complex queries needed to retrieve current vs. historical data.

❖ Business Benefits

- **Customer Analysis** – Track customer address or phone number changes for targeted marketing.
- **Price & Inventory Insights** – Analyze historical product pricing and stock variations.
- **Sales Trends** – Evaluate past product performance across different pricing strategies.

- ❖ The code snippet of the SCD implementation and the star schema design after the implementation are displayed below

3. Slowly Changing Dimensions (SCD)

❖ Customer Dimension (customer_dim_scd)

```
-- Creating a customer_dim SCD Table
CREATE TABLE customer_dim_scd (
    surrogate_key SERIAL PRIMARY KEY, -- Surrogate key for uniqueness
    customer_id VARCHAR(50),          -- Natural key from source system
    customer_name TEXT NOT NULL,
    address TEXT,
    telephone VARCHAR(20),
    effective_date DATE NOT NULL,      -- Date when this record became effective
    expiration_date DATE,              -- Date when this record was replaced
    current_flag BOOLEAN NOT NULL DEFAULT TRUE
);

-- Creating a Trigger Function for SCD Type 2 Updates
CREATE OR REPLACE FUNCTION scd_customer_dim_update()
RETURNS TRIGGER AS $$
BEGIN
    -- Insert OLD record into customer_dim_scd (archive old data)
    INSERT INTO customer_dim_scd (customer_id, customer_name, address, telephone, effective_date, expiration_date, current_flag)
    VALUES (OLD.customer_id, OLD.customer_name, OLD.address, OLD.telephone, CURRENT_DATE, CURRENT_DATE, FALSE);

    -- Allow the update to proceed normally in customer_dim
    RETURN NEW;
END;
$$ LANGUAGE plpgsql;

-- Attaching the trigger function to fire BEFORE updating customer_dim
CREATE TRIGGER trg_scd_customer_dim_update
BEFORE UPDATE ON customer_dim
FOR EACH ROW
EXECUTE FUNCTION scd_customer_dim_update();
```

3. Slowly Changing Dimensions (SCD)

◆ Product Dimension (product_dim_scd)

```
CREATE TABLE product_dim_scd (
    surrogate_key SERIAL PRIMARY KEY,      -- Unique identifier for each history record
    product_id VARCHAR(50),                -- Natural key (same for a given product)
    product_name TEXT NOT NULL,
    quantity_available INT,
    unit_price DECIMAL(10,2) NOT NULL,
    cost_price DECIMAL(10,2) NOT NULL,
    effective_date DATE NOT NULL,          -- When this version became effective
    expiration_date DATE,                  -- When this version was superseded (NULL if current)
    current_flag BOOLEAN NOT NULL DEFAULT TRUE -- Indicates if this is the current version in history
);

CREATE OR REPLACE FUNCTION scd_product_dim_capture_history()
RETURNS TRIGGER AS $$
BEGIN
    -- If an active history record exists for this product, mark it as expired.
    UPDATE product_dim_scd
    SET expiration_date = CURRENT_DATE,
        current_flag = FALSE
    WHERE product_id = OLD.product_id
        AND current_flag = TRUE;

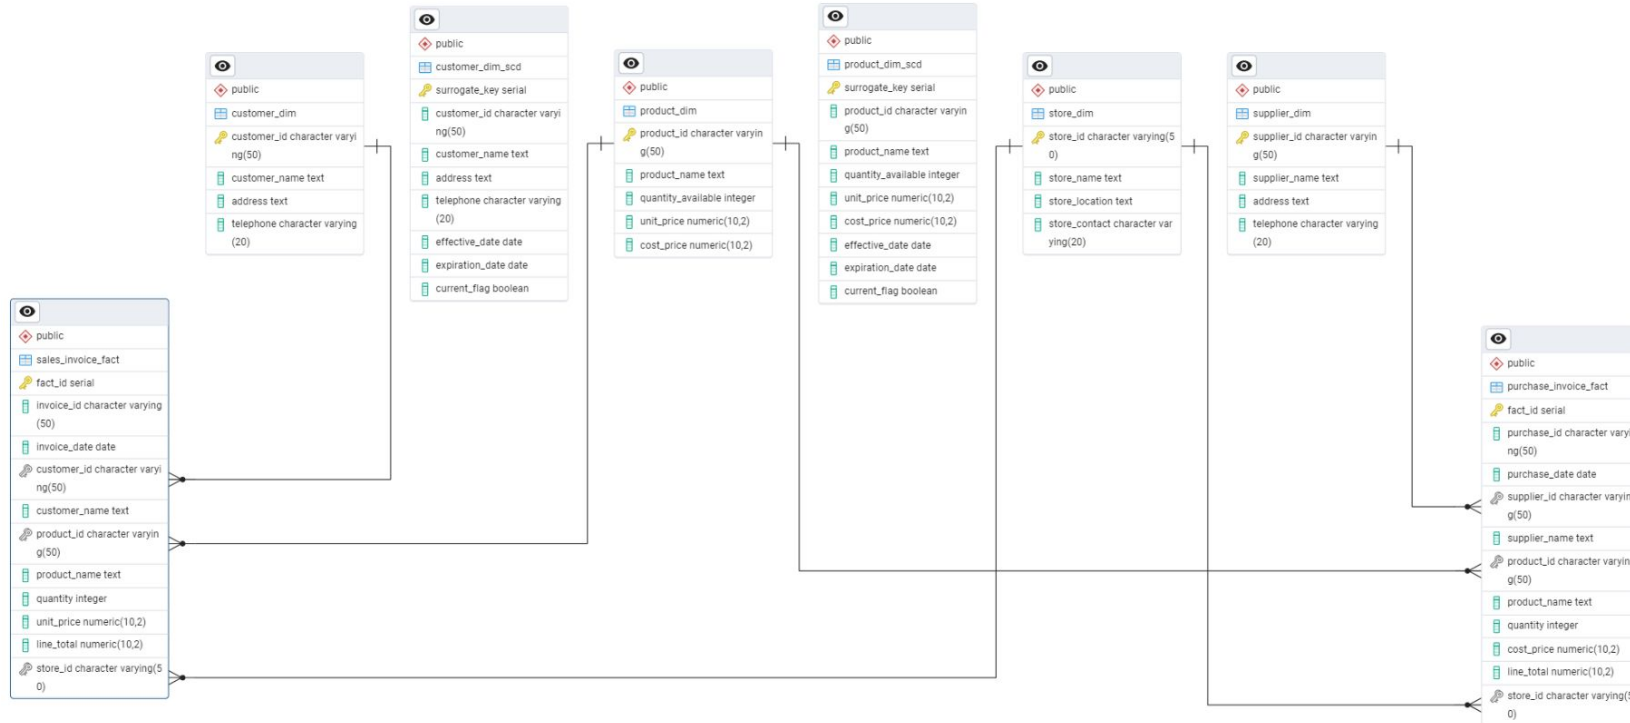
    -- Insert the old record into the history table.
    INSERT INTO product_dim_scd (
        product_id, product_name, quantity_available, unit_price, cost_price, effective_date, current_flag
    )
    VALUES (
        OLD.product_id, OLD.product_name, OLD.quantity_available, OLD.unit_price, OLD.cost_price, CURRENT_DATE, TRUE
    );

    RETURN NEW;
END;
$$ LANGUAGE plpgsql;

CREATE TRIGGER trg_capture_product_history
BEFORE UPDATE ON product_dim
FOR EACH ROW
EXECUTE FUNCTION scd_product_dim_capture_history();
```


3. Slowly Changing Dimensions (SCD)

❖ Star schema design after SCD implementation



4. ETL/ELT Strategy

❖ Data Integration Approach

- **Fact and Dimension Table Relationships** – Ensures proper linkage between sales, purchases, and product/customer/supplier details.
- **Trigger-Based Data Updates** – Automates data population and updates for real-time consistency.
- **Incremental Data Updates** – Uses AFTER INSERT and BEFORE INSERT triggers for efficient data changes.
- **Hybrid ETL/ELT Approach:**
 - ETL (Extract, Transform, Load) – Ensures data transformation before insertion.
 - ELT (Extract, Load, Transform) – Uses database triggers to transform after loading.

❖ Data Population & Automation

- **Product Quantity Management**
 - **On Purchase** → Increases quantity_available in product_dim
 - **On Sale** → Decreases quantity_available in product_dim.
- **Automated Data Enrichment**
 - **Sales Invoice** → Populates product_name and customer_name from respective dimensions.
 - **Purchase Invoice** → Populates product_name and supplier_name from respective dimensions.

4. ETL/ELT Strategy

❖ Data Quality & Consistency

- **Referential Integrity** – Ensures product_id, customer_id, and supplier_id exist before insertion.
- **Trigger-Based Consistency** – Ensures real-time updates across related tables.
- **Error Prevention** – Avoids missing or incorrect data by enforcing automatic lookups.
- **Timely Updates** – Data changes are reflected immediately after transactions.

❖ Performance Considerations

- **Trigger Optimization** – Prevents performance bottlenecks by limiting the scope of updates.
- **Indexing Key Columns** – Improves query efficiency on product_id, customer_id, and supplier_id.
- **Batch Processing for Large Volumes** – Reduces overhead when handling bulk inserts.

```
-- Trigger function: update product quantity on purchase
CREATE OR REPLACE FUNCTION update_quantity_on_purchase()
RETURNS TRIGGER AS $$
BEGIN
    UPDATE product_dim
    SET quantity_available = quantity_available + NEW.quantity
    WHERE product_id = NEW.product_id;
    RETURN NEW;
END;
$$ LANGUAGE plpgsql;

CREATE TRIGGER trg_update_quantity_purchase
AFTER INSERT ON purchase_invoice_fact
FOR EACH ROW
EXECUTE FUNCTION update_quantity_on_purchase();

-- Trigger function: update product quantity on sale
CREATE OR REPLACE FUNCTION update_quantity_on_sales()
RETURNS TRIGGER AS $$
BEGIN
    UPDATE product_dim
    SET quantity_available = quantity_available - NEW.quantity
    WHERE product_id = NEW.product_id;
    RETURN NEW;
END;
$$ LANGUAGE plpgsql;

CREATE TRIGGER trg_update_quantity_sales
AFTER INSERT ON sales_invoice_fact
FOR EACH ROW
EXECUTE FUNCTION update_quantity_on_sales();
```

4. ETL/ELT Strategy

```
-- Trigger function: populate product_name in sales_invoice_fact from product_dim
CREATE OR REPLACE FUNCTION populate_sales_invoice_product_name()
RETURNS TRIGGER AS $$
BEGIN
    SELECT p.product_name
    INTO NEW.product_name
    FROM product_dim p
    WHERE p.product_id = NEW.product_id;
    RETURN NEW;
END;
$$ LANGUAGE plpgsql;

CREATE TRIGGER trg_populate_sales_invoice_product_name
BEFORE INSERT ON sales_invoice_fact
FOR EACH ROW
EXECUTE FUNCTION populate_sales_invoice_product_name();

-- Trigger function: populate product_name and supplier_name in purchase_invoice_fact
CREATE OR REPLACE FUNCTION populate_purchase_invoice_names()
RETURNS TRIGGER AS $$
BEGIN
    -- Populate product_name from product_dim
    SELECT p.product_name
    INTO NEW.product_name
    FROM product_dim p
    WHERE p.product_id = NEW.product_id;

    -- Populate supplier_name from supplier_dim
    SELECT s.supplier_name
    INTO NEW.supplier_name
    FROM supplier_dim s
    WHERE s.supplier_id = NEW.supplier_id;

    RETURN NEW;
END;
$$ LANGUAGE plpgsql;
```

```
-- Populate supplier_name from supplier_dim
SELECT s.supplier_name
INTO NEW.supplier_name
FROM supplier_dim s
WHERE s.supplier_id = NEW.supplier_id;

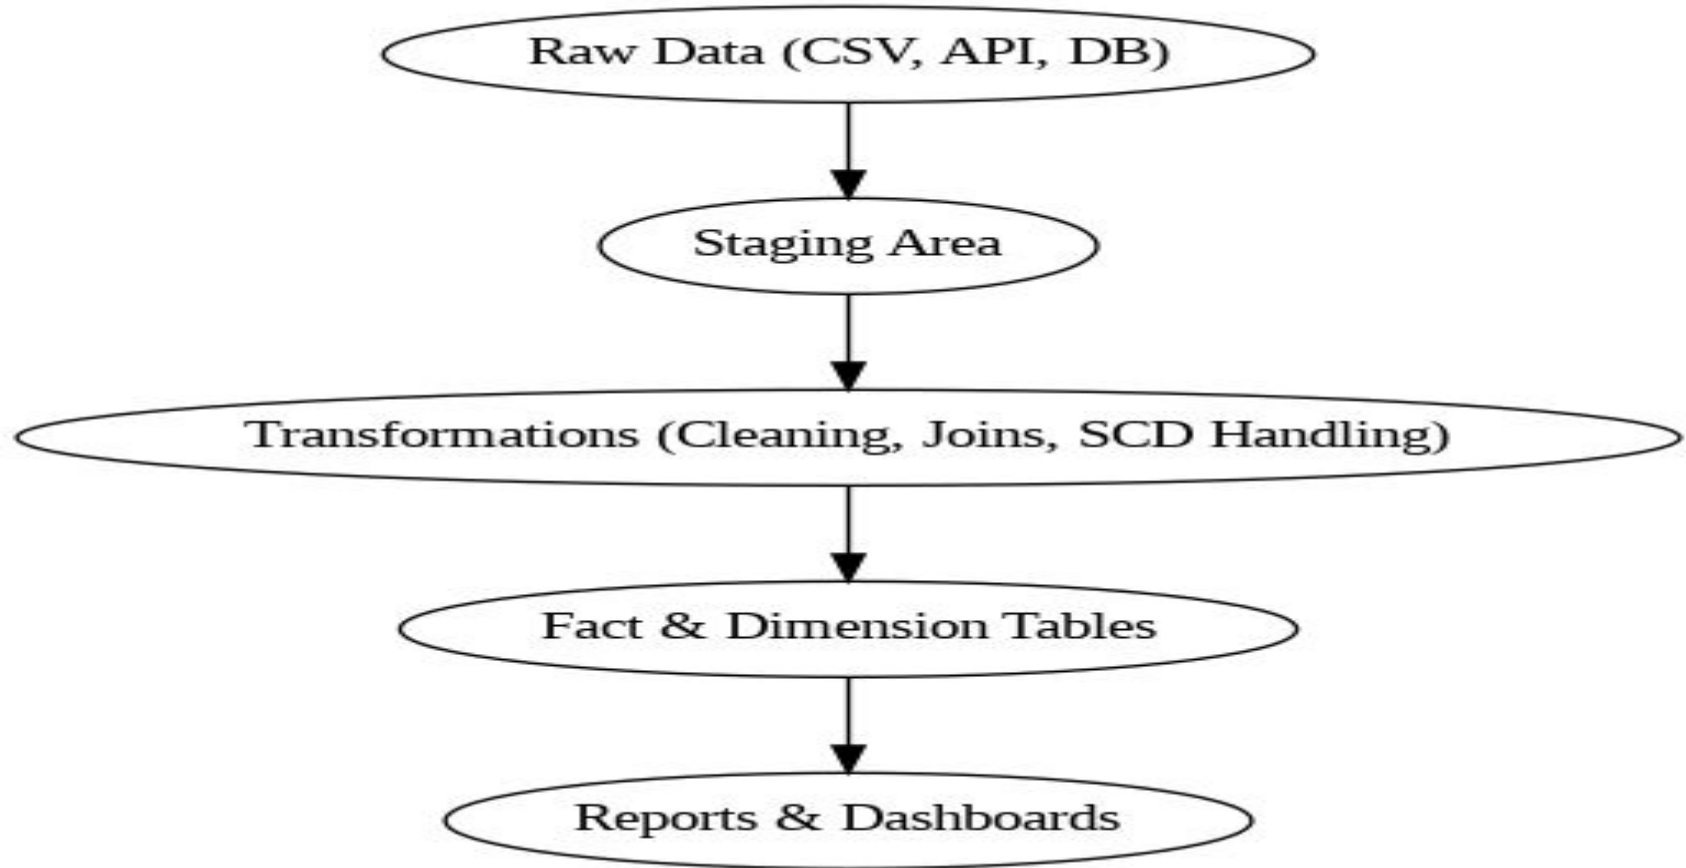
RETURN NEW;
END;
$$ LANGUAGE plpgsql;

CREATE TRIGGER trg_populate_purchase_invoice_names
BEFORE INSERT ON purchase_invoice_fact
FOR EACH ROW
EXECUTE FUNCTION populate_purchase_invoice_names();
```

```
-- Trigger function: populate customer_name in sales_invoice_fact from customer_dim
CREATE OR REPLACE FUNCTION populate_sales_invoice_customer_name()
RETURNS TRIGGER AS $$
BEGIN
    SELECT c.customer_name
    INTO NEW.customer_name
    FROM customer_dim c
    WHERE c.customer_id = NEW.customer_id;
    RETURN NEW;
END;
$$ LANGUAGE plpgsql;

CREATE TRIGGER trg_populate_sales_invoice_customer_name
BEFORE INSERT ON sales_invoice_fact
FOR EACH ROW
EXECUTE FUNCTION populate_sales_invoice_customer_name();
```

4. ETL/ELT Strategy



5. Advanced Analysis & Scalability

❖ Hierarchical Relationships

➤ Multi-Level Drill-Down Analysis

- Organize products by category, subcategory, and brand to enhance reporting.
- Enable time-based analysis (Year → Quarter → Month → Day) for trend evaluation

➤ Parent-Child Relationships.

- Support roll-up and drill-down queries for regional performance tracking.

➤ Aggregated Reporting

- Design summary tables for quick retrieval of high-level insights.
- Use materialized views to precompute common aggregations.

❖ Performance Optimization

➤ Database Indexing

- Use composite indexes for frequent filter combinations (e.g., date & product).

➤ Partitioning Strategy

- Partition fact tables by date range to speed up historical data retrieval.

➤ Surrogate Key Implementation

- Replace natural keys with integer surrogate keys for faster joins.

➤ Query Optimization Techniques

- Use caching for frequently accessed data to reduce database load.

➤ Scalability & Growth Considerations

- Design schema to accommodate future data growth with minimal restructuring

Relevance to Retailers

1. Supports Analytical Reporting
 2. Enhances Drill-Down Capabilities
 3. Ensures Historical Tracking
 4. Improves Scalability & Performance
 5. Facilitates Business Decision-Making
-



END OF REPORT